

Algorithms and Satisfiability

Lecture 4: External-Memory Algorithms and Data Structures

DAT6 spring 2023
Simonas Šaltenis



AALBORG UNIVERSITY
DENMARK

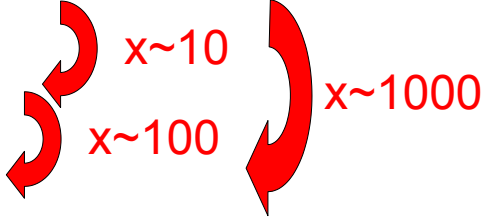
External Mem. Algorithms and DS



- Goals of the lecture:
 - *to understand the external memory model and the principles of analysis of algorithms and data structures in this model;*
 - *to understand the algorithms of B-tree and its variants and to be able to analyze them;*
 - *to understand the main principles of external tree structures;*
 - *to understand how the different versions of **merge-sort** derived algorithms work in external memory;*
 - *to understand why the amount of available main-memory is an important parameter for the efficiency of external-memory algorithms.*

Memory hierarchy, prices




- In 2021, people created ~2.5 exabytes (million TBs) *per day!*
 - Where do we store that data?
- Prices:
 - HDD price: ~0.02 \$/GB
 - SSD price: ~0.15 \$/GB
 - DRAM price: ~5-10 \$/GB
- *Memory-hierarchy is still very relevant in the age of big data!*

Sources: <https://techjury.net/>, <https://pcpartpicker.com/>

Hard disk I



- In *real systems*, we need to cope with data that does not fit in main memory
- Reading a data element from the hard-disk:
 - *Seek* with the head
 - *Wait* while the necessary sector rotates under the head
 - *Transfer* the data



Hard disk II



- Modern hard drives:
 - *Seek time*: 4ms-10ms
 - *Spindle speed*: ~10K RPM \Rightarrow *Half of rotation*: ~3ms
 - *Transfer rate*: 500 MB/s \Rightarrow *Transferring 1 byte*: 0.000003ms
- Conclusions:
 1. It makes sense to *read and write in large blocks* – disk pages (4 – 32Kb)
 2. *Sequential* access is much faster than *random* access
 3. Disk access is much slower than main-memory access

SSDs, Memory Hierarchy



- The same, although to less extent is true for *flash-based solid state drives* (SSDs):
 - Efficient to read/write (especially write) in larger blocks
 - Sequential/random I/O difference is less pronounced than in disks.
- Depth of the memory hierarchy (access latency):
 - **DRAM**(~50ns) – **x4000** → **SSD**(~0.2ms) – **x50** → **HDD**(10ms)
If = 1s, then > 1 hour, > 2 days
- Memory hierarchy consisting of several levels of **CPU caches** and **DRAM**:
 - Again, data between levels is transferred in *blocks*
 - In contrast to disk drives and SSDs, block reads and writes are not explicit – controlled by hardware/low level system software

External memory model



- Running time: in *page accesses* or “I/Os”
- *B* – page size is an important parameter:
 - Not “just” a constant:
 - ♦ $\Theta(\log_2 n) \neq \Theta(\log_B n)$
 - ♦ $\Theta(N) \neq \Theta(N/B)$
 - ♦ *Example*: $N = 256\text{MB} / 8 \text{ bytes_per_object}$;
 $B = 4\text{KB} / 8 \text{ bytes_per_object}$; 0.1 ms disk access
 - ▲ N disk accesses = 3200s = 53 minutes
 - ▲ N/B disk accesses = 6.4s
- Operations:
 - ♦ `DiskRead(x:pointer_to_a_page)`
 - ♦ `DiskWrite(x:pointer_to_a_page)`
 - ♦ `AllocatePage():pointer_to_a_page`

Writing algorithms



- The typical working pattern for algorithms:

```
01 ...  
02  $x \leftarrow$  a pointer to some object  
03 DiskRead(x)  
04 operations that access and/or modify x  
05 DiskWrite(x) //omitted if nothing changed  
06 other operations, only access no modify  
07 ...
```

- Pointers in data-structures point to disk-pages, not locations in memory

“Porting” main-memory DSs




- Why not “just” use the main-memory data structures and algorithms in external memory?
- Consider a balanced binary search tree.
 - $A, B, C, D, E, F, G, H, I$
- Options:
 - Each node gets a separate disk page – waist of space and search is just $\Theta(\log_2 n)$
 - Nodes are somehow packed to make disk pages full – search may still be $\Theta(\log_2 n)$ in the worst-case

B-trees




- We are concerned only with keys
- The nodes have high *fan-out* (many children) = $\Theta(B)$
 - *Degree* of a tree t :
 - ♦ $Min_fan-out = t$, $Max_fan-out = 2*t = B / index_entry_size$
 - Root is the exception: can have as little as two children
- B-tree is a balanced tree, and all leaves have *the same depth*: $h = \Theta(\log_t n) = \Theta(\log_B n)$



B-trees, nodes



- Internal nodes
 - $t - 1$ to $2t - 1$ keys
 - $\text{pointer}_1 \text{ key}_1 \text{ pointer}_2 \text{ key}_2 \text{ pointer}_3 \text{ key}_3 \dots \text{pointer}_x \text{ key}_x \text{ pointer}_{x+1}$
 - $\text{key}_1 \leq \text{key}_2 \leq \text{key}_3 \leq \dots \leq \text{key}_x$
 - For the first and last pointers: $\text{pointer}_1.\text{key} \leq \text{key}_1$
 - ...and $\text{key}_x < \text{pointer}_{x+1}.\text{key}$
 - For the remaining pointers: $\text{key}_{i-1} < \text{pointer}_i.\text{key} \leq \text{key}_i$
- Leave nodes
 - Only keys.



Searching on B-trees



- The root node is normally “always” in main memory.
 - No need to perform a DiskRead on the root.
- Search is very similar to a search in a binary search tree
 - Instead of making a binary branching decision at each node, we make a $(j+1)$ -way branching decision, where j is the number of keys in a node.

Pseudo code



- x is a node and $x.n$ is the number of keys in the node.
- k is the key that we are searching for.
- $x.key_i$ is the i -th key of node x ; and $x.c_i$ is the i -th pointer of node x .

B-TREE-SEARCH(x, k)

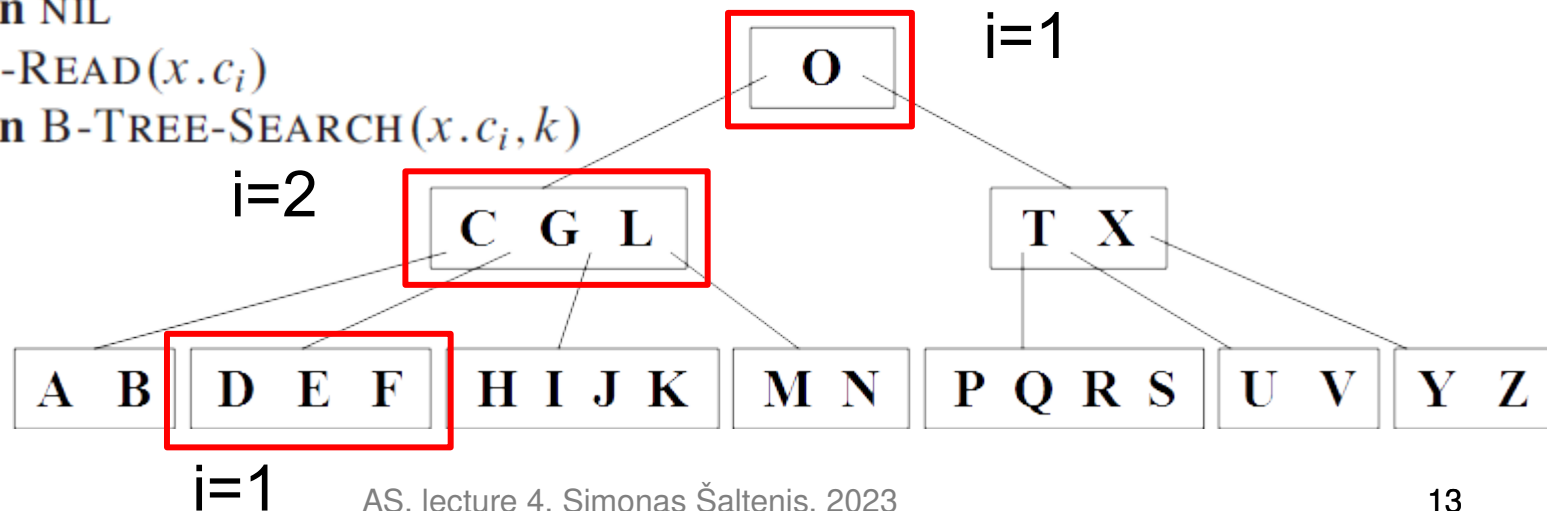
```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return ( $x, i$ )
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

Searching for “D”, i.e., $k = D$

B-Tree-Search(root, D)

Disk access: $O(h) = O(\log_t n)$


CPU: $O(t h) = O(t \log_t n)$



B⁺-trees




- B⁺-trees is a variant of B-trees:
 - All data keys are in leaf nodes
 - ◆ What is the height?
 - Leaf-nodes are connected into a (doubly) linked list
 - Search is very similar to a search in a binary search tree
 - ◆ Always goes to a leaf
 - ◆ Range searches are convenient
 - ◆ Cost: $\Theta(\log_B n + k/B)$



B⁺-trees: Insertion




- Skeleton of the algorithm:
 - *Down-phase*: recursively traverse down and find the leaf (as in search)
 - *Up-phase*: Insert the key. If necessary, *split* nodes and propagate the splits up the tree
- Assumption:
 - In the *down-phase* pointers to traversed nodes are saved in the stack as there are *no parent pointers*!
- Insert M:




B⁺-trees: Node Splitting



- Leaf node (*copy* the middle key to the parent)



- Internal node (*move* the middle key to the parent, as in B-tree)




- The tree grows when the root is split into two nodes and their parent becomes the new root.

B⁺-trees: Example



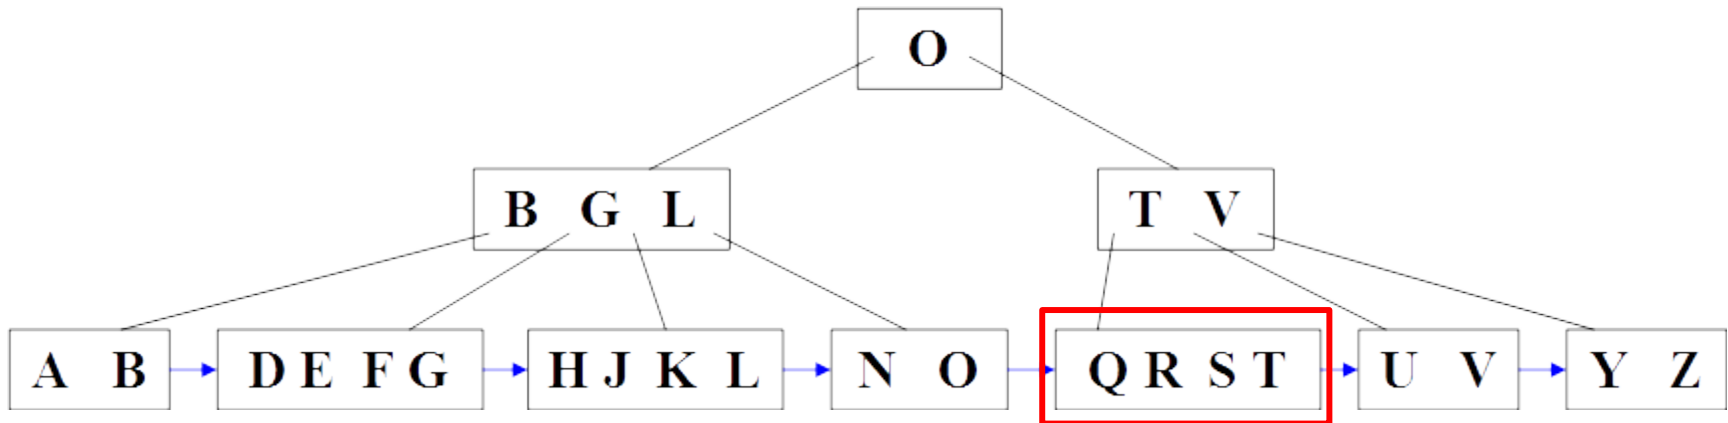
- *Insert P* (maximum fan-out = 5 children = 4 data entries)
- *What is the cost?*
 - $\Theta(\log_B n)$
- *How much memory is used?*
 - $\Theta(\log_B n)$ – can be reduced to $\Theta(1)$.




Example



- *Insert P* (assume that at most 4 keys)




Down-phase: Leaf node
(**copy** the middle key to the parent)



B⁺-trees: insert exercise



- Exercise:
 - Assume maximum fan-out = 5 children = 4 data entries
 - *Insert I, C*




- *So, why parent pointers are usually not used in B-trees, in contrast to binary search trees?*


Example




- At most 4 keys
- After inserting I, C
- How does the root node look like?



Example




- Insert I
 - HJKL becomes HIJKL. Then split HIJ, KL. J is copied to its parent.



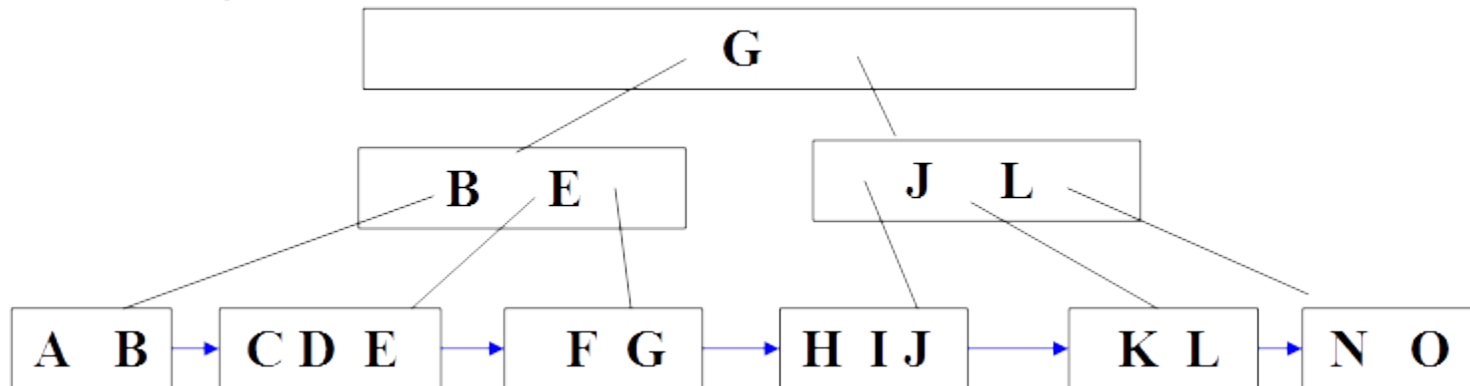
Example



- Insert C
 - DEFG becomes CDEFG. Split CDE, FG. E is copied to its parent.



- BEGJL splits to BE and JL and a new root with G is created.



B⁺-trees: Deletion




- Opposite of insertion:
 - Phase 1: traverse down to find the key in a leaf
 - Phase 2: remove the key and traverse up handling underfull nodes
- Tree shrinking: if the root has only one child, remove the root.


B⁺-trees: Deletion (internal nodes)



- Underfull handling, case 1: *Distributing*



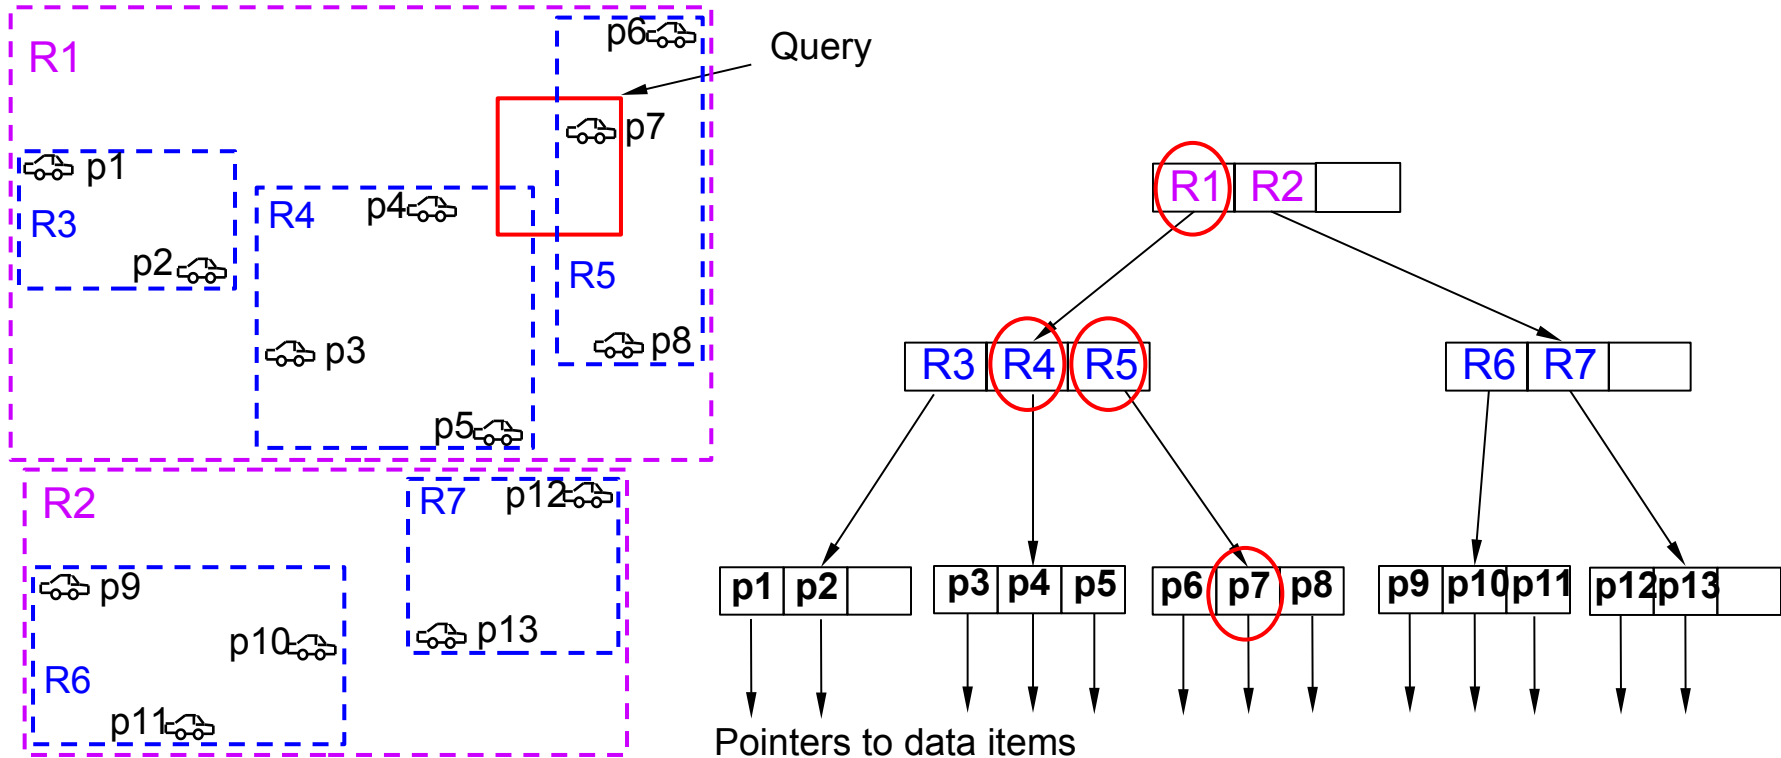
- Underfull handling, case 2: *Merging*



R-trees



- Example



Grow-Post Trees




Grow-Post trees (Generalized Search Trees - GiST)

Bounding predicate (BP) = *something that describes entries in a subtree*

Building blocks of algorithms:

- **Consistent**(BP, Q) – returns *true* if results of query Q can be under BP (in the R-tree, MBR intersects Q).
- **Penalty**(BP, E) – returns an estimate how “worse” BP becomes if E is inserted under it.
- **Union**($node$) – computes a BP of a collection of entries (in the R-tree, computes an MBR, minimum bounding rectangle)
- **PickSplit**($node$) – splits a page of entries into two groups



External DS: Summary



- Two practical data structures (n is the number of pages):
 - **B-trees**: supports point and range queries, insertions, deletions
 - ♦ Point query: $\Theta(\log_B n)$
 - ♦ Range query: $\Theta(\log_B n + k/B)$
 - ♦ Insertion, deletion: $\Theta(\log_B n)$
 - **R-trees**: supports multi-dimensional point and range queries, on point and extended objects:
 - ♦ Point/range query, deletion: $\Theta(n)$, but usually much better on average
 - ♦ Insertion: $\Theta(\log_B n)$
 - Both structures have $\Theta(n)$ size.

External-Memory Sorting



- External-memory algorithms
 - When data do not fit in main-memory
- External-memory sorting
 - Rough idea: sort pieces that fit in main-memory and “merge” them
- Main-memory merge sort:
 - The main part of the algorithm is Merge
 - Let's merge:
 - ♦ 3, 6, 7, 11, 13
 - ♦ 1, 5, 8, 9, 10

Main-Memory Merge Sort



Merge-Sort (A)

01 **if** length(A) > 1 **then**

02 Copy the first half of A into array A1

03 Copy the second half of A into array A2

04 **Merge-Sort** (A1)


05 **Merge-Sort** (A2)

06 **Merge** (A, A1, A2)

} *Divide*
} *Conquer*
} *Combine*

- Running time?

Merge-Sort Recursion Tree

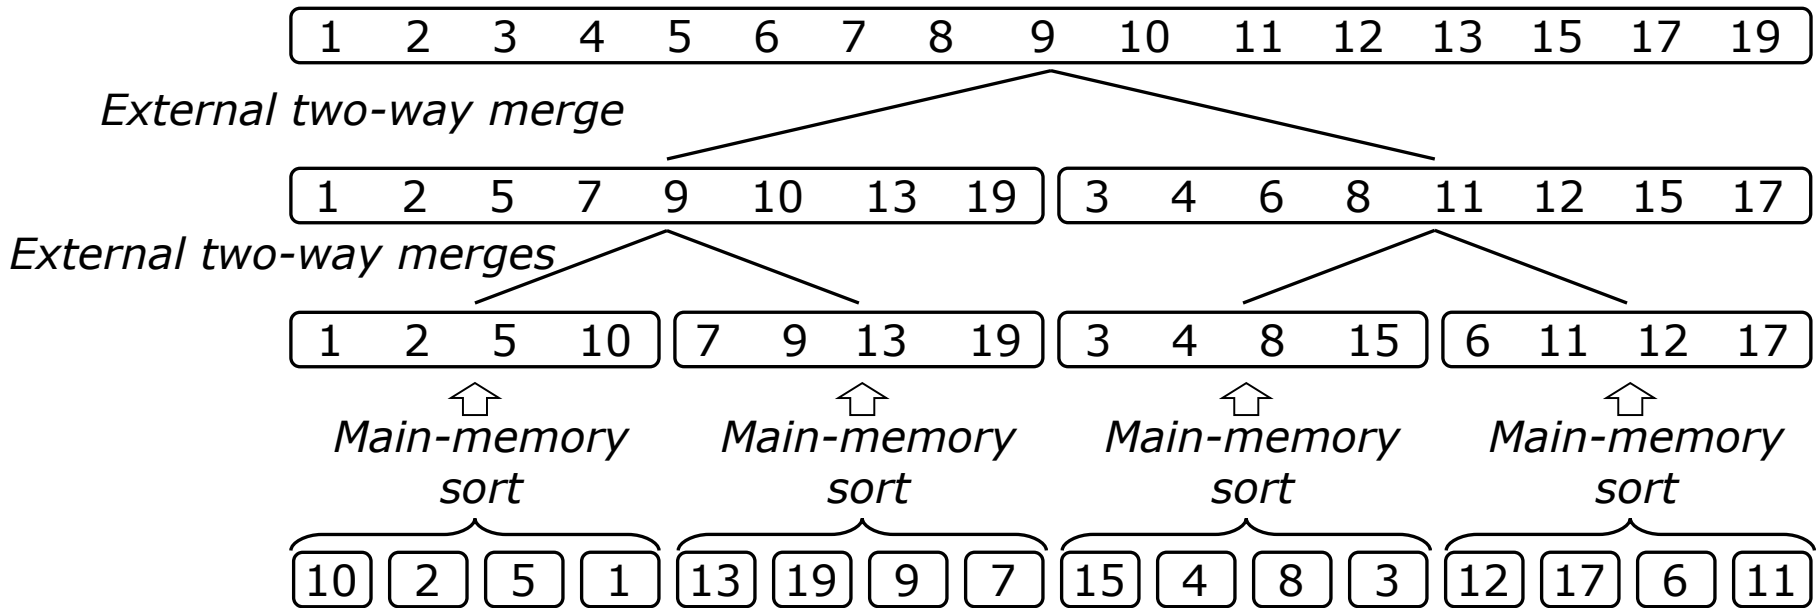


- In each level: merge *runs* (sorted sequences) of size x into runs of size $2x$, decrease the number of runs twofold.
- What would it mean to run this on a file in external memory?

External-Memory Merge-Sort



- Idea: increase the size of initial runs!
 - Initial runs – the size of available main memory (M data elements)



External-Memory Merge Sort

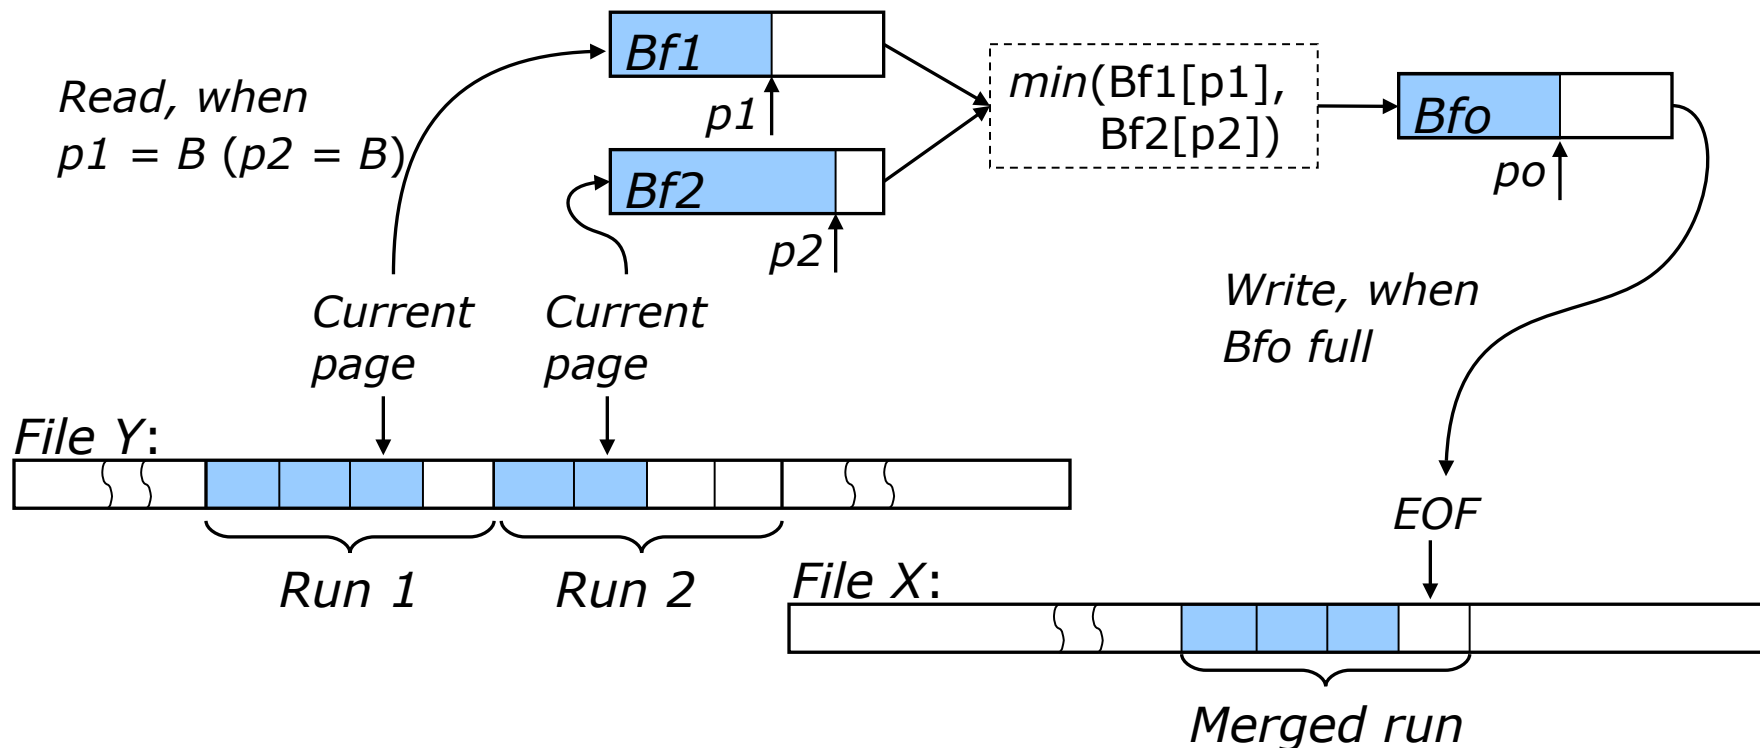


- Input file X , empty file Y
- *Phase 1*: Repeat until the end of file X :
 - Read the next M elements from X
 - Sort them in main-memory
 - Write them at the end of file Y
- *Phase 2*: Repeat while there is more than one run in Y :
 - Empty X
 - *MergeAllRuns*(Y, X)
 - X is now called Y , Y is now called X

External-Memory Merging



- *MergeAllRuns*(Y, X): repeat until the end of Y :
 - Call *TwowayMerge* to merge the next two runs from Y into one run, which is written at the end of X
- *TwowayMerge*: uses three main-memory arrays of size B

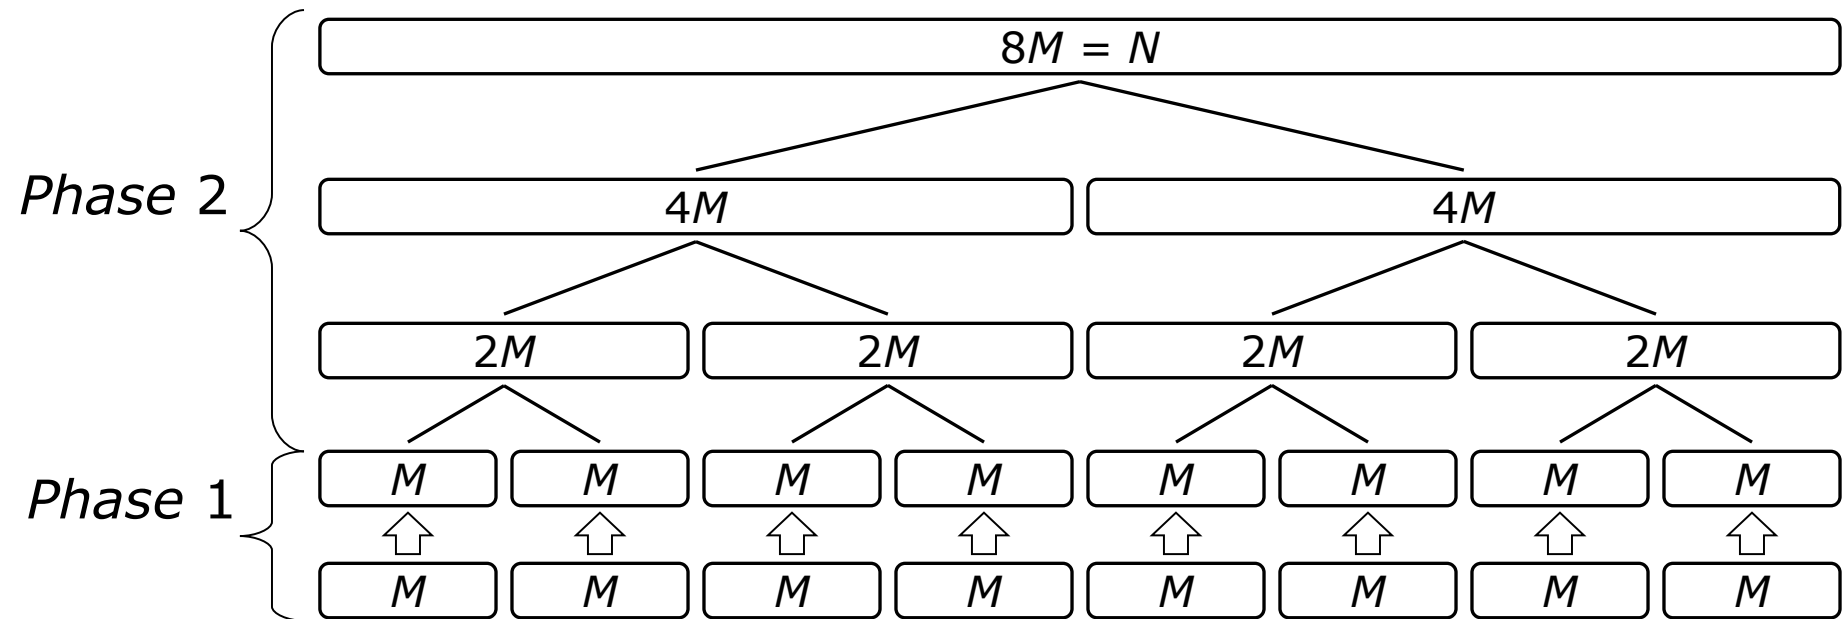


Analysis: Assumptions



- Assumptions and notation:
 - Disk page size:
 - ♦ B data elements
 - Data file size:
 - ♦ N elements, $n = N/B$ disk pages
 - Available main memory:
 - ♦ M elements, $m = M/B$ pages

Analysis



- Phase 1:
 - Read file X, write file Y: $2n = \Theta(n)$ I/Os
- Phase 2:
 - One iteration: Read file Y, write file X: $2n = \Theta(n)$ I/Os
 - Number of iterations: $\log_2 N/M = \log_2 n/m$

Analysis: Conclusions




- Total running time of external-memory merge sort: $\Theta(n \log_2 n/m)$
- We can do better!
- Observation:
 - Phase 1 uses all available memory
 - Phase 2 uses just 3 pages out of m available!!!

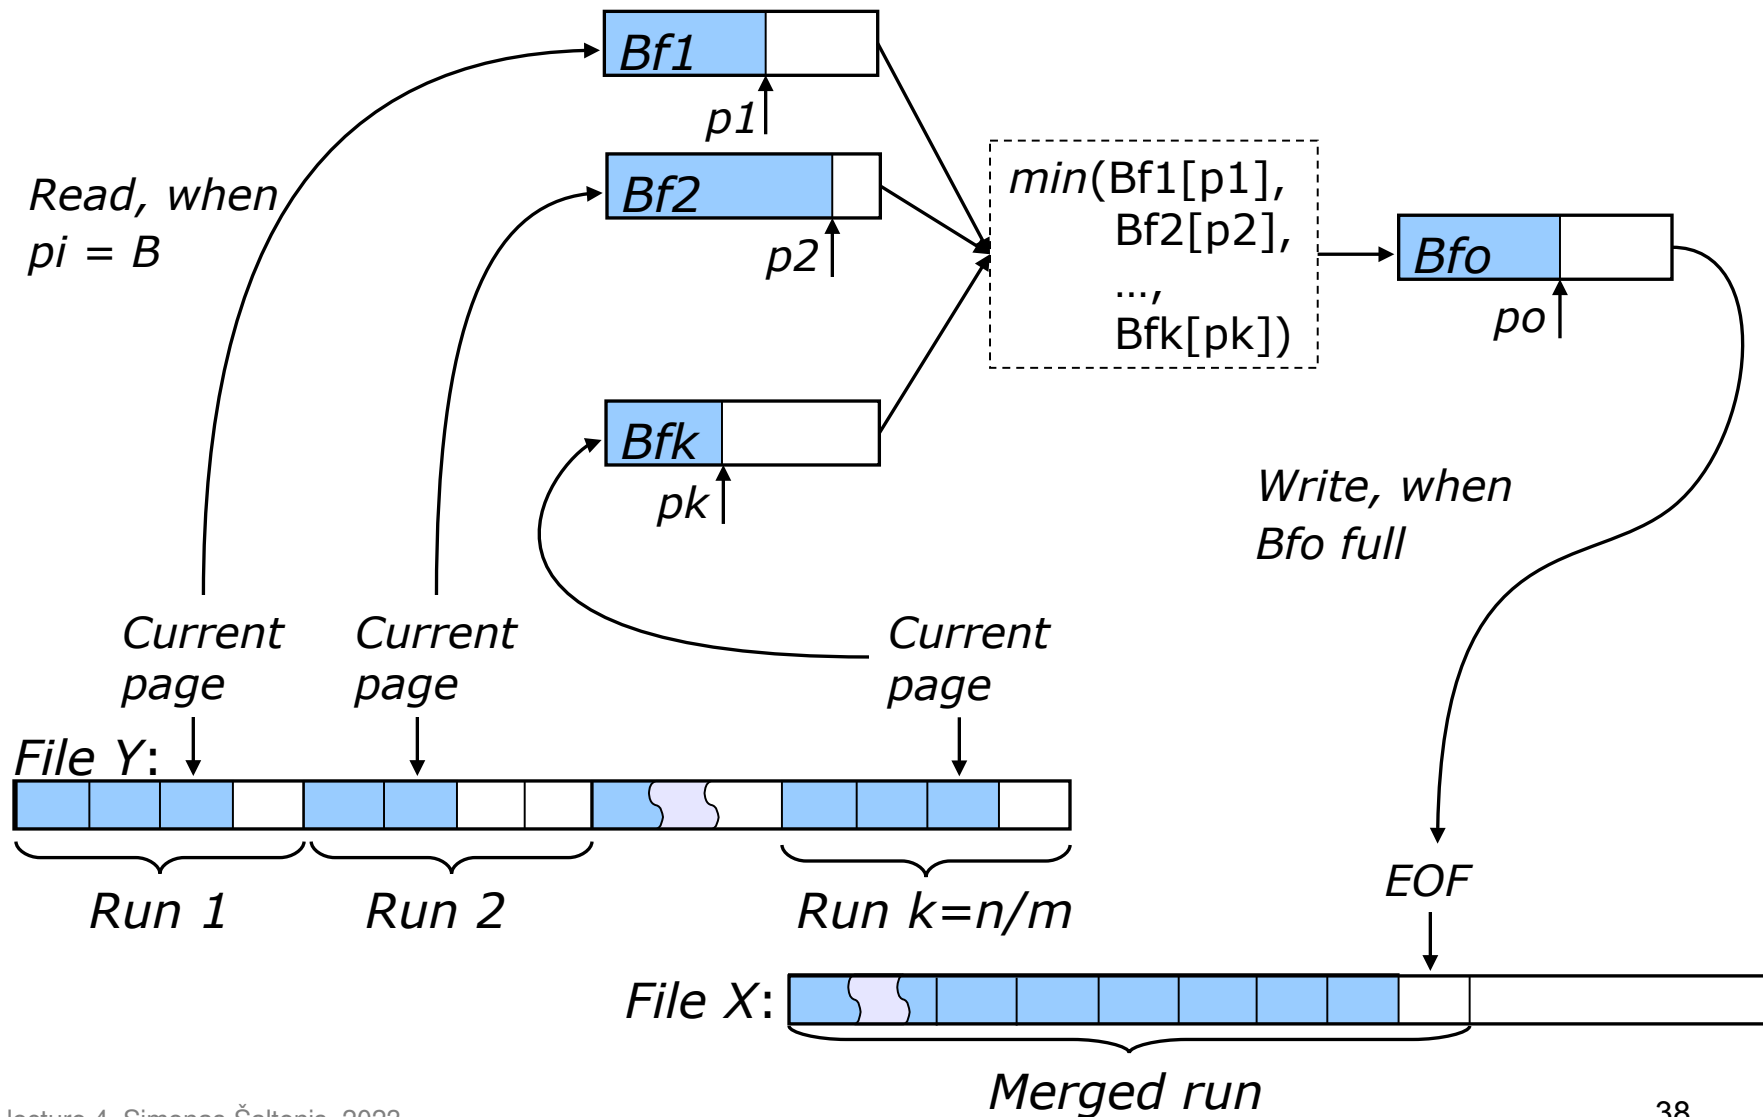
Two-Phase, Multiway Merge Sort



- Idea: merge all runs at once!
 - Phase 1: the same (do internal sorts)
 - Phase 2: perform *MultiwayMerge*(Y, X)



Multiway Merging



Analysis of TPMMS



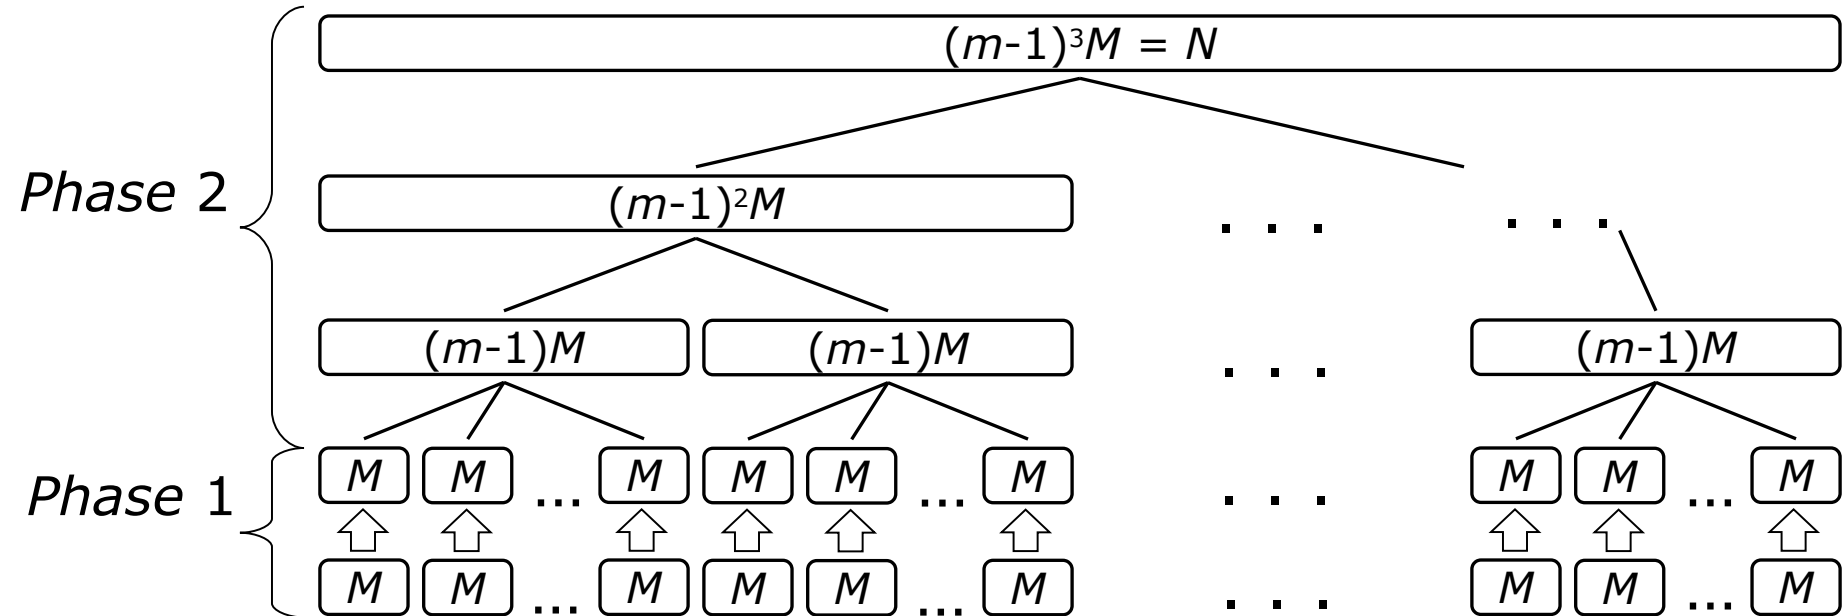
- Phase 1: $\Theta(n)$, Phase 2: $\Theta(n)$
- Total: $\Theta(n)$ I/Os!
- The catch: files only of “limited” size can be sorted
 - Phase 2 can merge a maximum of $m-1$ runs.
- Which means: $N/M \leq m-1$ ($n/m \leq m-1$)
 - *How large files can we sort with TPMMS on a machine with 128MiB main memory and disk page size of 16KiB?*

General Multiway Merge Sort



- What if a file is **very** large or memory is small?
- General *multiway merge sort*:
 - Phase 1: the same (do internal sorts)
 - Phase 2: do as many iterations of merging as necessary until only one run remains
 - ♦ Each iteration repeatedly calls *MultiwayMerge*(Y , X) to merge groups of $m-1$ runs until the end of file Y is reached

Analysis



- Phase 1: $\Theta(n)$, each iteration of phase 2: $\Theta(n)$
- How many iterations are there in phase 2?
 - Number of iterations: $\log_{m-1} N/M = \Theta(\log_m n)$
- Total running time: $\Theta(n \log_m n)$ I/Os

Conclusions



- External sorting can be done in $\Theta(n \log_m n)$ I/O operations for any n
 - This is asymptotically optimal
- In practice, we can usually sort in $\Theta(n)$ I/Os
 - Use two-phase, multiway merge-sort