

Algorithms and Satisfiability

Lecture 5: Parallel Algorithms

DAT6 spring 2023
Simonas Šaltenis



AALBORG UNIVERSITY
DENMARK

Parallel algorithms

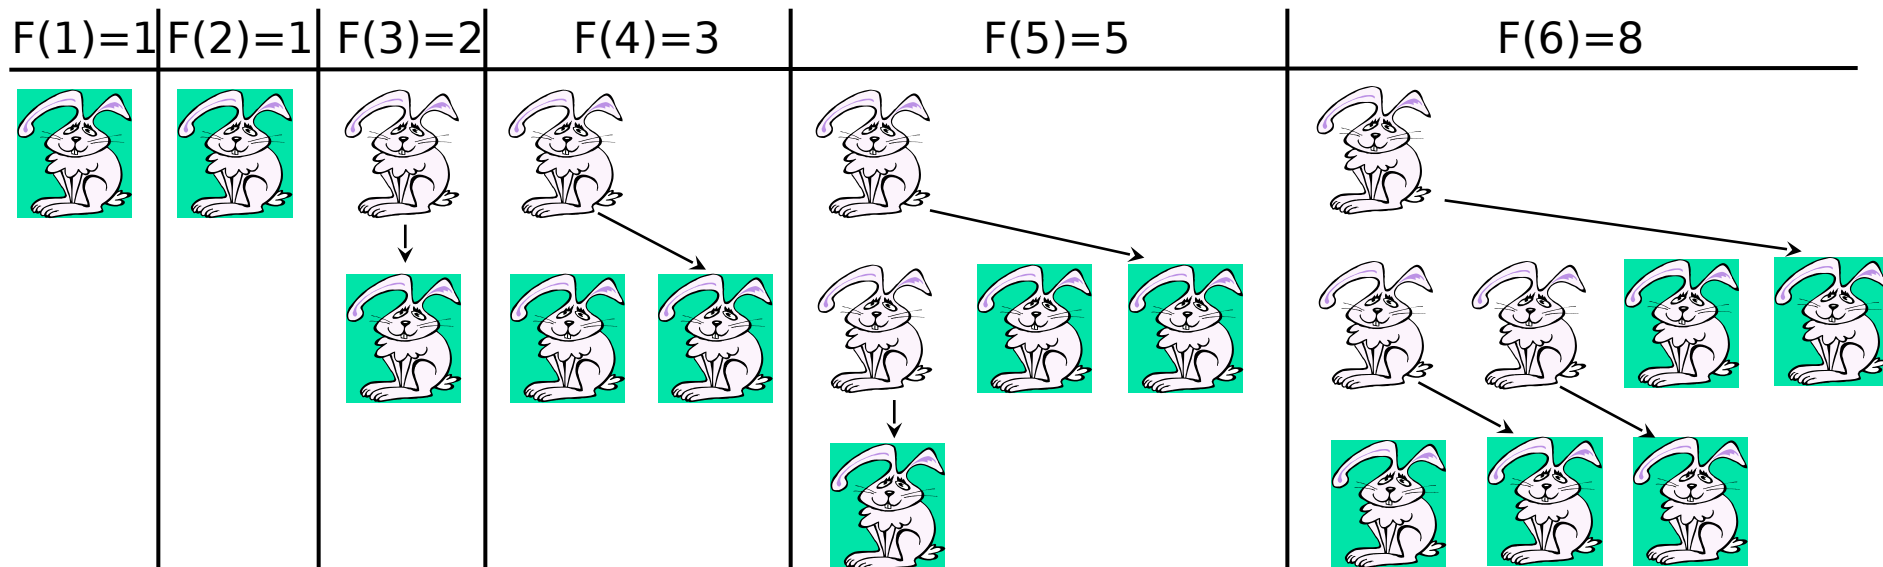


- Goals of the lecture:
 - *to understand the model of **dynamic multithreading** (aka **fork-join parallelism**);*
 - *to understand **work**, **span**, and **parallelism** — the concepts necessary for the analysis of parallel algorithms;*
 - *to understand and be able to analyze the **parallel merge sort** algorithm.*

Fibonacci Numbers



- *Leonardo Fibonacci (1202):*
 - A rabbit starts producing offspring on the second generation after its birth and produces one child each generation
 - How many rabbits will there be after n generations?



Fibonacci Numbers (2)

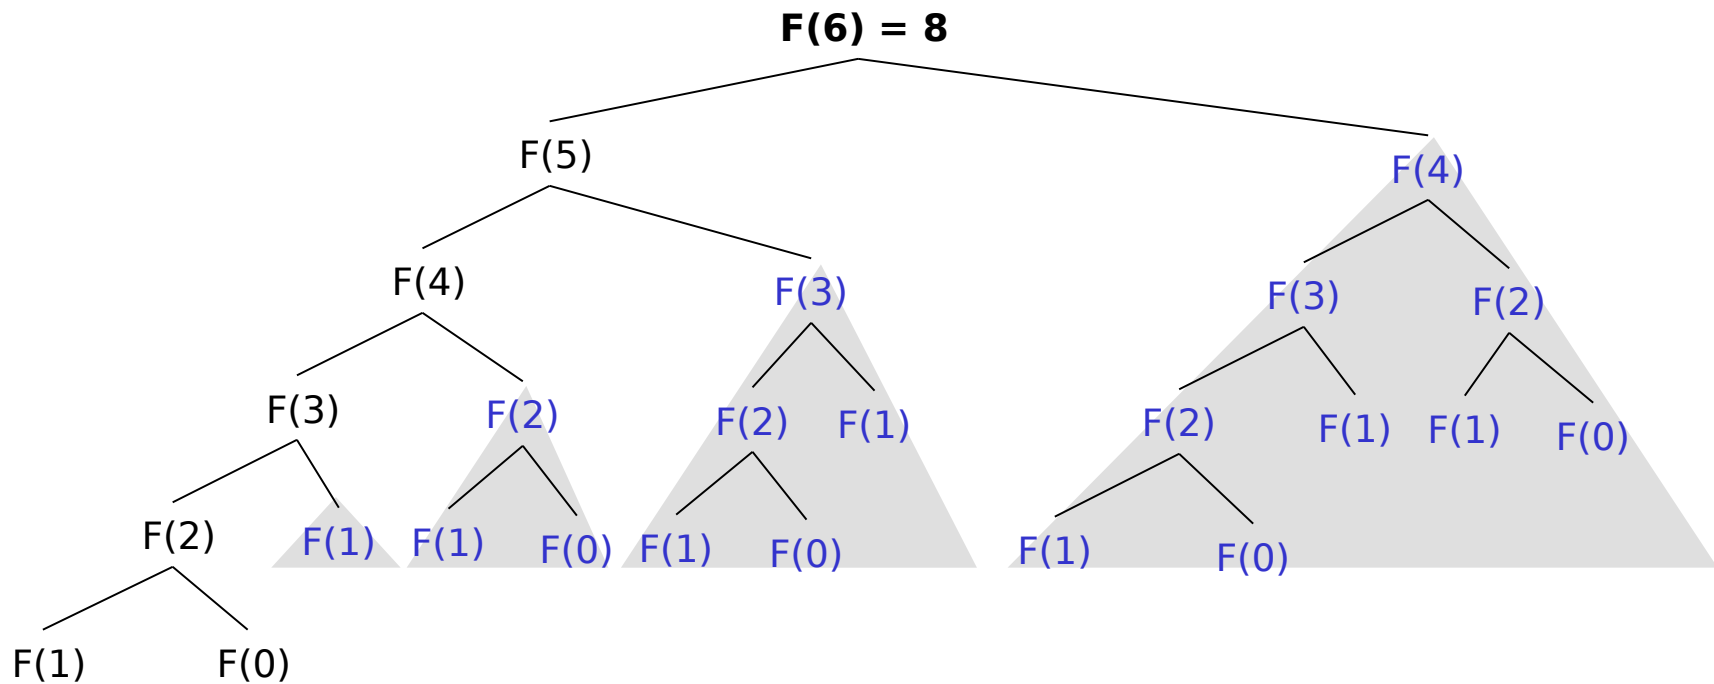


- $F(n) = F(n-1) + F(n-2)$
- $F(0) = 0, F(1) = 1$
 - 0, 1, 1, 2, 3, 5, 8, 13, 21, 34 ...

```
Fibonacci (n)
01 if n ≤ 1 then return n
02 else
03   x = Fibonacci(n-1)
04   y = Fibonacci(n-2)
05 return x + y
```

- Straightforward recursive procedure is slow!
- Why? How slow?
- Let's draw the recursion tree

Fibonacci Numbers (3)



Fibonacci Numbers (4)



- How many summations are there $W(n)$?
 - $W(n) = W(n - 1) + W(n - 2) + 1$
 - $W(n) \geq 2W(n - 2) + 1$ and $W(1) = W(0) = 0$
 - Solving the recurrence we get
$$W(n) \geq 2^{(n-1)/2} - 1 \approx 1.4^{n-1}$$
 - Precisely $W(n) = \Theta(\varphi^n)$, where φ is the *golden ratio* $(1+\sqrt{5})/2 \approx 1.618$
- Running time is **exponential**.

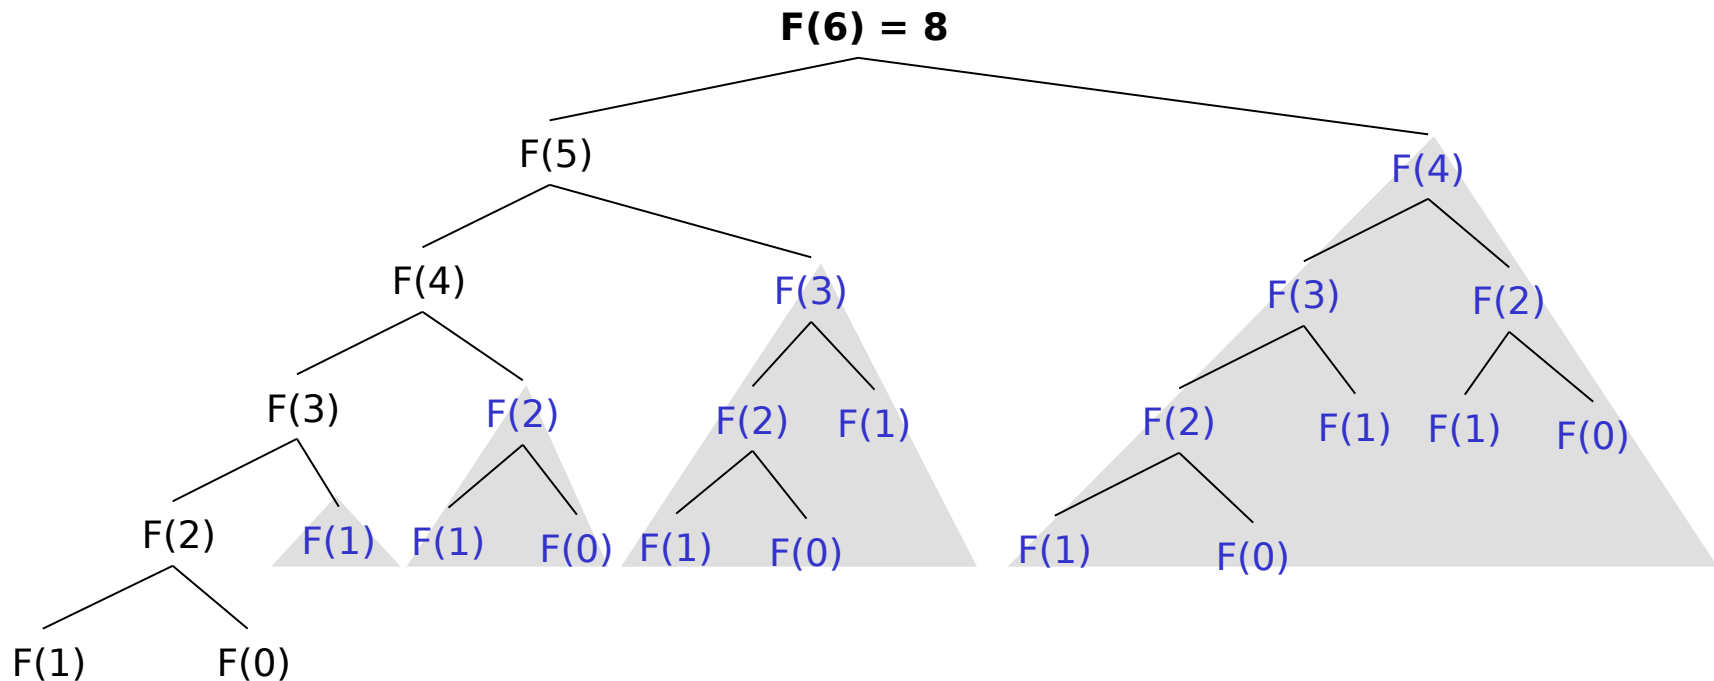
Multithreaded version



- What if we can do the two recursive calls in parallel
 - Using the so-called *nested parallelism*

```
FibonacciP(n)
01 if n ≤ 1 then return n
02 else
03   x = spawn FibonacciP(n-1)
04   y = FibonacciP(n-2)
05   sync
06 return x + y
```

FibonacciP analysis



- “Running time”:
 - $S(n) = \max(S(n-1), S(n-2)) + 1 = S(n-1) + 1$
 - Thus $S(n) = \Theta(n)$

Work, Span, Parallelism



- Three main concepts (informally):
 - **Work**: the running time on a machine with one-processor (T_1).
 - ♦ Fibonacci: $\Theta(\varphi^n)$
 - **Span**: the running time on a machine with infinite processors (T_∞).
 - ♦ Fibonacci: $\Theta(n)$
 - **Parallelism** = $Work/Span$ – how many processors on average are used by the algorithm.
 - ♦ Fibonacci: $\Theta(\varphi^n / n)$
- More formally:
 - **Computation log** – a DAG of serial strands of instructions (vertices) and dependencies (edges) between them.
 - Work = the number of vertices in the computation log.
 - Span = the length of the *longest path (critical path)* in the computation log.

Computation DAG



- Using an example of computing Fibonacci number of 4.

```
FibonacciP(n)
```

```
01 if n ≤ 1 then return n
```

```
02 else
```

```
03   x = spawn FibonacciP(n-1)
```

```
04   y = FibonacciP(n-2)
```

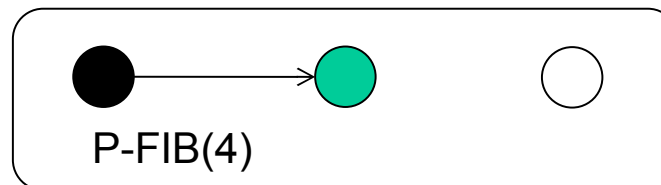
```
05   sync
```

```
06 return x + y
```

● Lines 1-3

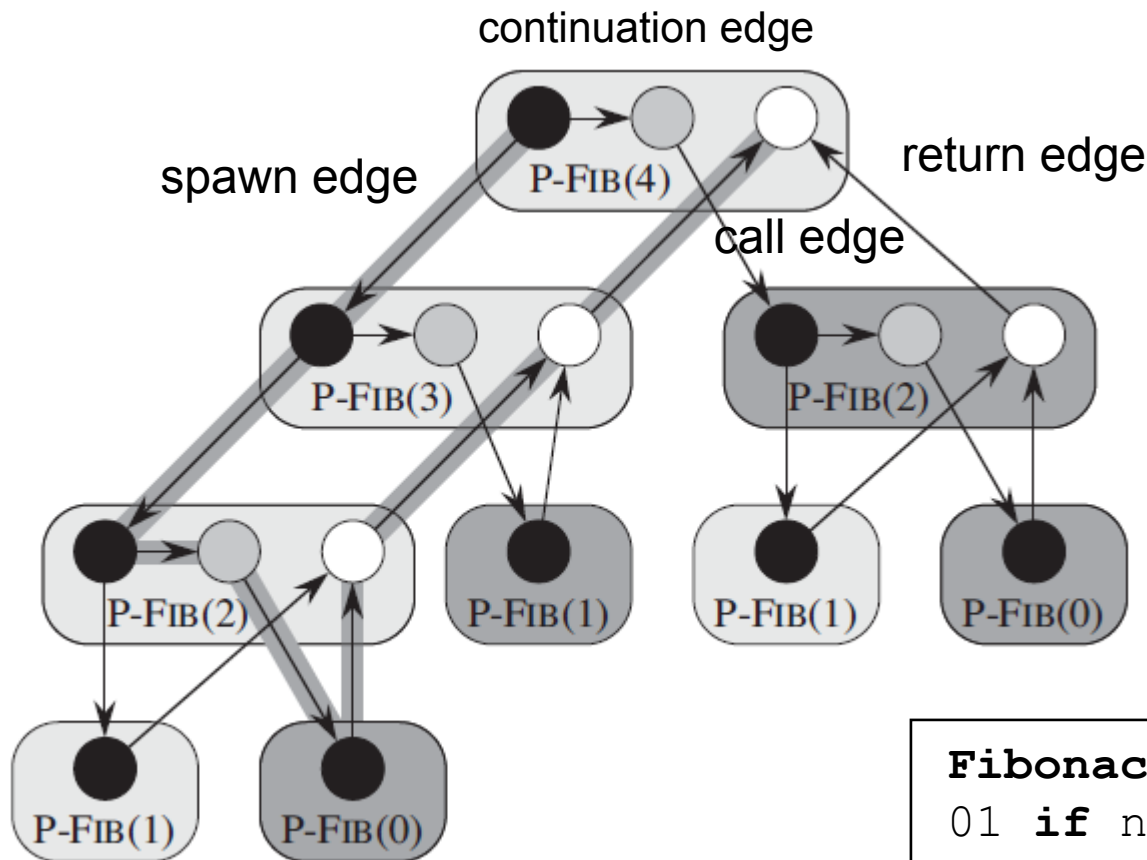
● Lines 4-5

○ Line 6

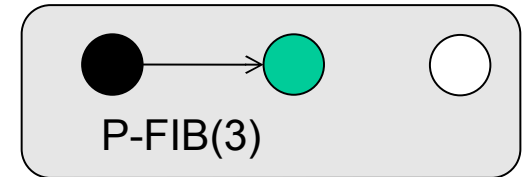


Computation DAG

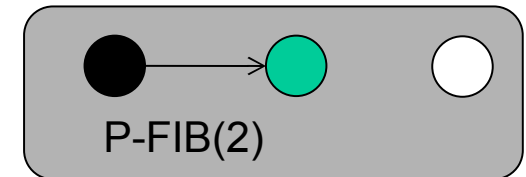
Edge(u, v) means that u must execute before v .



Spawned procedure



Called procedure



Work: number of vertices, 17
Span: the length of the longest path (critical path), 8

FibonacciP(n)

```
01 if  $n \leq 1$  then return  $n$ 
02 else
03    $x = \text{spawn FibonacciP}(n-1)$ 
04    $y = \text{FibonacciP}(n-2)$ 
05   sync
06 return  $x + y$ 
```

Work law and span law



- Notation
 - Work T_1
 - Span T_∞
 - Multithreaded computation on P processors: T_P
- Work law: $T_P \geq T_1 / P$
 - An ideal parallel computer with P processors can do at most P units of work.
- Span law: $T_P \geq T_\infty$
 - An ideal parallel computer with P processors cannot run any faster than a machine with unlimited number of processors.

Assumptions



- The *fork-join parallelism (dynamic multithreading)* environment:
 - Shared-memory multi-core system
 - *Concurrency platform – task-parallel programming* :
 - ♦ Takes care of allocating work to physical threads (in other words: scheduling logical threads on physical threads)
 - ♦ Takes care of synchronization, consistent access to memory
 - Pseudocode keywords: *spawn*, *sync* and *parallel*
 - ♦ Indicates *potential* (or *logical*) parallelism: what *may* run in parallel.
 - We do not consider locking, race conditions, etc:
 - ♦ Parallel threads are *independent* – they work on separated items of data.
 - We abstract from actual physical scheduling:
 - ♦ It can be shown that simple greedy scheduling works well enough.

Speedup, Slackness



- When running on an actual system with P physical threads:
 - *Slackness* of a computation: $Parallelism / P$.
 - *What does it mean when slackness < 1 ? Slackness > 1 ?*
 - $Speedup = T_1 / T_P$.
 - *Perfect linear speedup*, when speedup = P .

Mini quiz



- Considering the case for computing P-Fib(4).
- We already know that the work $T_1 = 17$ and the span $T_\infty = 8$
- Consider the following setups, each setup corresponds to a machine with P processing units. Which one is the most likely setup to achieve the *perfect linear speedup*?
- A: $P=2$, B: $P=3$, C: $P=4$?
- Go to [Socrative](#) and vote.

To Summarize

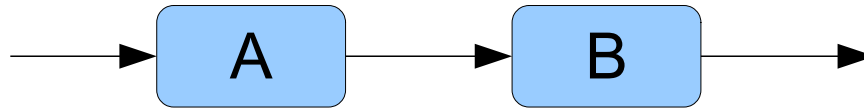


Notation	Meaning
T_1	Work, the running time on a machine with one processor.
T_∞	Span, the running time on a machine with infinite processors.
T_P	The running time on a machine with P processors.
$T_P \geq T_1 / P$	Work law
$T_P \geq T_\infty$	Span law
T_1 / T_P	Speedup. Speedup must be $\leq P$ according to the work law. When speedup is equal to P , it achieves perfect speed up .
T_1 / T_∞	Parallelism. The maximum possible speedup that can be achieved on any number of processors
T_1 / PT_∞	Slackness = Parallelism/ P . The larger the slackness, the more likely to achieve perfect speed up. When slackness is less than 1, it is impossible to achieve perfect speed up.

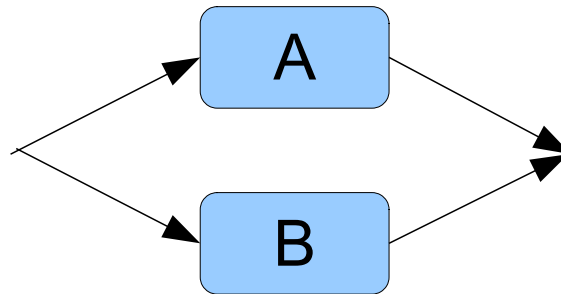
Computing span



- Sequential execution:
 - Work and span: $T(A \text{ followed_by } B) = T(A) + T(B)$



- Parallel execution:
 - Work: $T_1(A \text{ in_parallel_with } B) = T_1(A) + T_1(B)$
 - Span: $T_\infty(A \text{ in_parallel_with } B) = \max(T_\infty(A), T_\infty(B))$



Goal of algorithm design



- Goal of the parallel algorithm design – *increase parallelism*.
 - Usually achieved by decreasing span (remember, parallelism = W/S)
 - It may pay off to slightly increase work, if span can be decreased significantly (in practice, relevant for highly parallel systems, such as supercomputers, but also GPUs)
- Intra-operation parallelism vs. inter-operation parallelism.

Side Note: Efficient Fibonacci



- The efficient $O(n)$ serial algorithm:
 - Simple application of “*dynamic programming*” (or *memoized evaluation* of the recursive version)

```
FibonacciImproved(n)
01 if  $n \leq 1$  then return n
02 Fim2  $\leftarrow$  0
03 Fim1  $\leftarrow$  1
04 for i  $\leftarrow$  2 to n do
05   Fi  $\leftarrow$  Fim1 + Fim2
06   Fim2  $\leftarrow$  Fim1
07   Fim1  $\leftarrow$  Fi
05 return Fi
```

- Can be actually done in $O(\lg n)$ additions and multiplications.

Parallel loops



- Denoted by the **parallel** keyword

```
ArrayCopy(A, B)
01 parallel for i = 1 to sizeof(A) do
02     B[i] = A[i]
```

- Analysis of span:
 - $S(n) = O(\lg n) + \max_i S_{\text{iteration}}(i)$
 - *Why?*
 - Parallel loop is implemented by divide-and-conquer

```
ArrayCopyRecursive(A, B, l, r)
01 if l > r then return
02 if l = r then B[l] = A[r]
03 else
04     q =  $\lceil (l+r)/2 \rceil$ 
05     spawn ArrayCopyRecursive(A, B, l, q-1)
06         ArrayCopyRecursive(A, B, q, r)
```

Examples



- Exchanging neighboring elements:

ArrayExchange (A)

```
01 parallel for i = 1 to |sizeof(A)/2| do  
02     tmp          = A[i*2]  
03     A[i*2]       = A[i*2-1]  
04     A[i*2-1]     = tmp
```



- Compute the largest stock price difference :

LargestSpike (A)

```
01 lspike = 0  
02 parallel for i = 1 to sizeof(A)-1 do  
03     if |A[i+1] - A[i]| > lspike then  
04         lspike = |A[i+1] - A[i]|  
05 return lspike
```



Merge Sort



Merge-Sort(A, p, r)

01 **if** p < r **then**

02 q = $\lfloor (p+r)/2 \rfloor$

03 **Merge-Sort**(A, p, q)

04 **Merge-Sort**(A, q+1, r)

05 **Merge**(A, p, q, r)

} *Divide*
} *Conquer*
} *Combine*

- Running time?

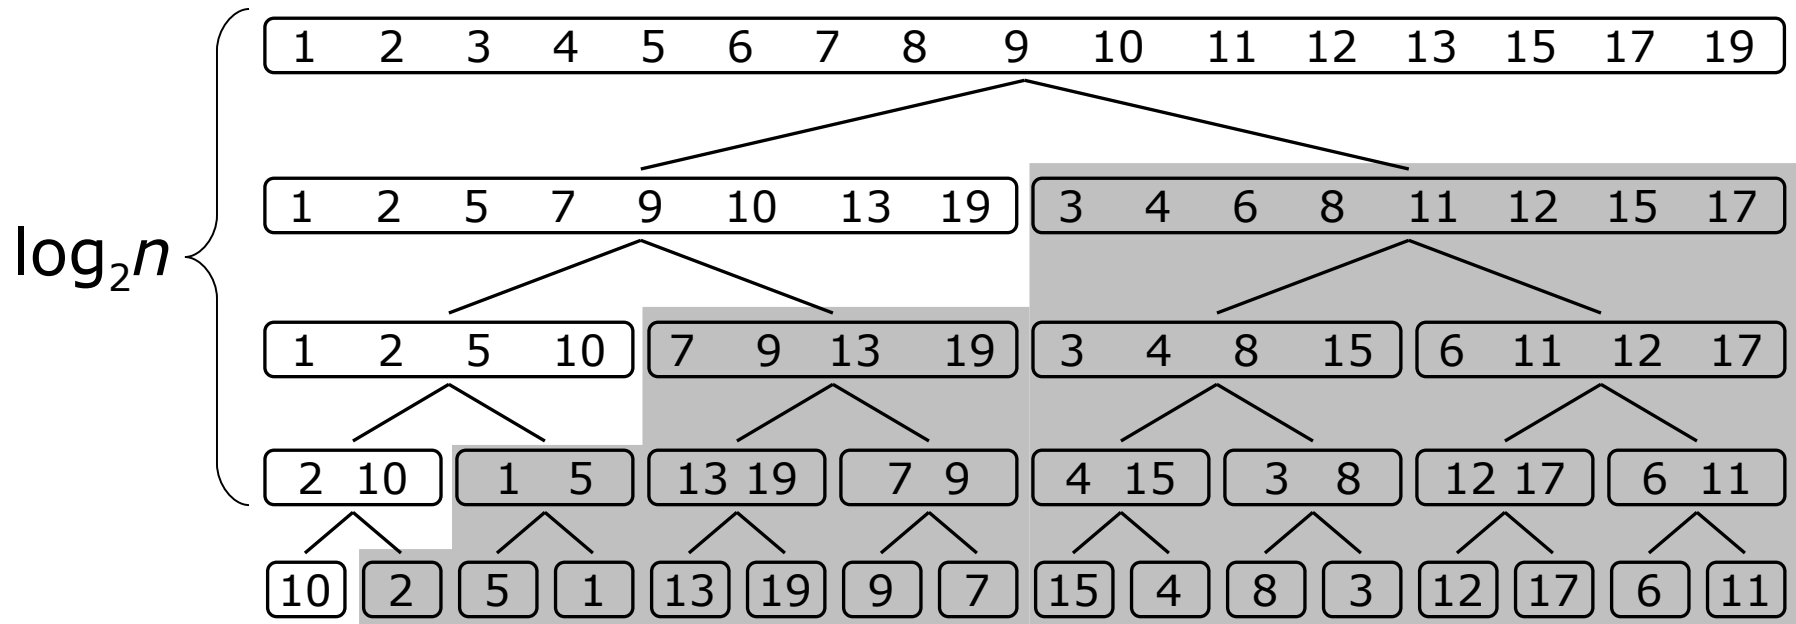
Parallelising Merge Sort



```
Merge-Sort' (A, p, r)
01 if p < r then
02     q = ⌊(p+r)/2⌋
03     spawn Merge-Sort' (A, p, q)
04     Merge-Sort' (A, q+1, r)
05     sync
06     Merge (A, p, q, r)
```

- Work
 - $W(n) = 2W(n/2) + \Theta(n)$
 - $W(n) = \Theta(n \lg n)$
- Span:
 - $S(n) = S(n/2) + \Theta(n)$
 - $S(n) = \Theta(n)$
- Parallelism: $W(n)/S(n) = \Theta(\lg n)$. Rather low...

Merge-Sort Recursion Tree

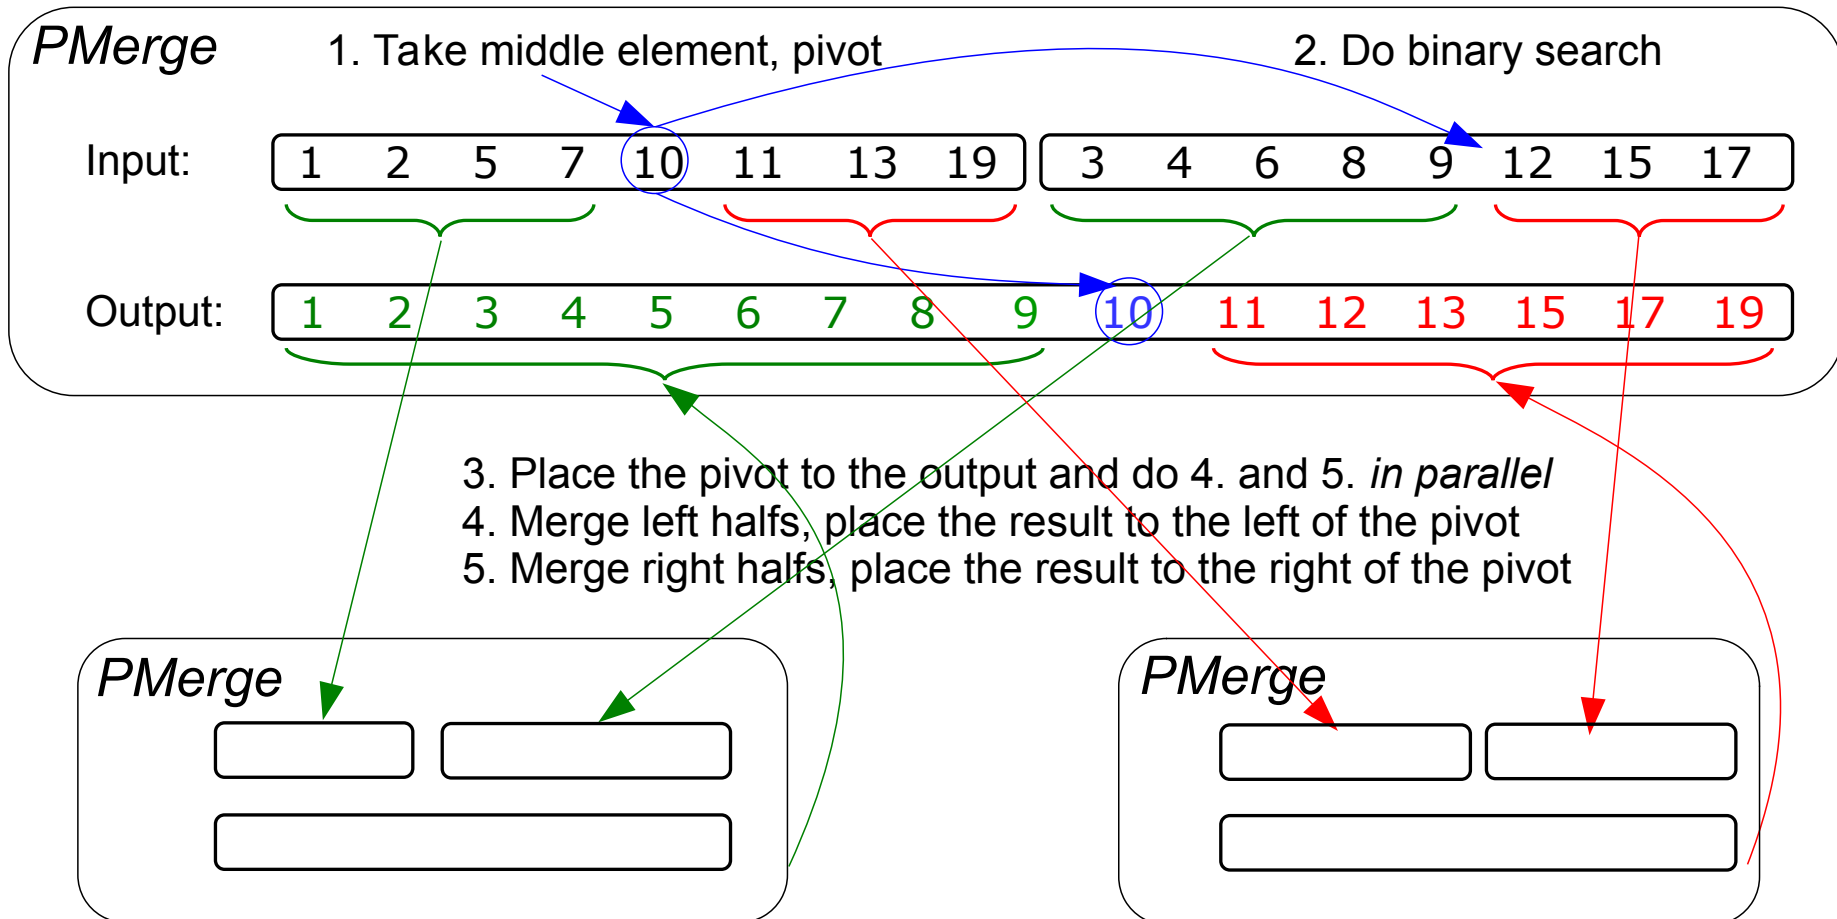


- Problem – merging is very serial:
 - At the top level, only one processor does $\Theta(n)$ work in serial!
 - At the second level, only two processors do $\Theta(n)$ work.
 - ...

Multithreaded merging



- Main idea – make the algorithm divide-and-conquer and use nested parallelism.



Multithreaded merging analysis



- One key idea: do binary search in the smaller of the two arrays!
 - This ensures that the *largest* of the two recursive calls works with *at most* $3n/4$ elements, where n = the sum of sizes of the two arrays.
 - *Why?*
- Span:
 - $S(n) = S(3n/4) + \Theta(\lg n)$
 - *What is the solution?*
 - $S(n) = \Theta(\lg^2 n)$
- Work:
 - Can be shown to be $\Theta(n)$.

Multithreaded merge sort



```
PMerge-Sort(A, p, r)
01 if p < r then
02     q = ⌊(p+r)/2⌋
03     spawn PMerge-Sort(A, p, q)
04     PMerge-Sort(A, q+1, r)
05     sync
06     PMerge(A, p, q, r)
```

- Work:
 - The same recurrence and solution: $W(n) = \Theta(n \lg n)$
- Span:
 - $S(n) = S(n/2) + \Theta(\lg^2 n)$
 - $S(n) = \Theta(\lg^3 n)$
- Parallelism:
 - $W(n) / S(n) = \Theta(n \lg n) / \Theta(\lg^3 n) = \Theta(n / \lg^2 n)$

To summarize



- Work: $\Theta(n \lg n)$

	Merge Procedure	Span	Parallelism
Naïve merge	$\Theta(n)$	$\Theta(n)$	$\Theta(\lg n)$
P-Merge	$\Theta(\lg^2 n)$	$\Theta(\lg^3 n)$	$\Theta(n / \lg^2 n)$

Goal of the multi-threaded algorithm design



- Goal of the multi-threaded algorithm design – *increase parallelism*.
 - $Parallelism = work / span$
 - Usually achieved by decreasing span
 - ♦ MergeSort without P-Merge and with P-Merge
 - ♦ $\Theta(n)$ vs. $\Theta(\lg^3 n)$
 - It may pay off to slightly increase work, if span can be decreased significantly (in practice, relevant for highly parallel systems, such as supercomputers, GPUs)